

```
In [1]: 1 import numpy as np
        2 import pandas as pd
```

```
In [2]: 1 data=pd.read_csv(r"uber.csv")
        2 # test_df=pd.read_csv(r"test.csv")
        3 print (data.shape)
        4 print (data.columns)
```

```
(200000, 9)
Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
      'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
      'dropoff_latitude', 'passenger_count'],
      dtype='object')
```

```
In [3]: 1 data_x = data.iloc[:,0:-1].values
        2 data_y = data.iloc[:, -1].values
        3 print(data_y)
```

```
[1 1 1 ... 2 1 1]
```

```
In [5]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null  int64
1   key                   200000 non-null  object
2   fare_amount          200000 non-null  float64
3   pickup_datetime      200000 non-null  object
4   pickup_longitude     200000 non-null  float64
5   pickup_latitude      200000 non-null  float64
6   dropoff_longitude     199999 non-null  float64
7   dropoff_latitude     199999 non-null  float64
8   passenger_count      200000 non-null  int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
In [6]: 1 data["pickup_datetime"]=pd.to_datetime(data['pickup_datetime'])
        2
```

In [7]:

```
1 data.head()
```

Out[7]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.73
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.72
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.72
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.75
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.72

As this is Taxi fare data and

we know there are many factors which affect the price of taxi like Travelled distance Time of Travel Demand and Availability of Taxi Some special places are more costlier like Airport or other places where there might be toll

In [8]:

```
1 data.describe()
```

Out[8]:

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	2.000000e+05	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	2.771250e+07	11.359955	-72.527638	39.935885	-72.525292	39.935885
std	1.601382e+07	9.901776	11.437787	7.720539	13.117408	7.720539
min	1.000000e+00	-52.000000	-1340.648410	-74.015515	-3356.666300	-74.015515
25%	1.382535e+07	6.000000	-73.992065	40.734796	-73.991407	40.734796
50%	2.774550e+07	8.500000	-73.981823	40.752592	-73.980093	40.752592
75%	4.155530e+07	12.500000	-73.967154	40.767158	-73.963658	40.767158
max	5.542357e+07	499.000000	57.418457	1644.421482	1153.572603	1644.421482

Here first thing which we can see is minimum value of fare is negative which is -62 which is not the valid value, so we need to remove the fare which are negative values. Secondly, passenger_count minimum value is 0 and maximum value is 208 which impossible, so we need to remove them as well, for safer side we can think that a taxi can have maximum 7 people.

```
In [9]: 1 #Lets check if there is any null value
        2 data.isnull().sum()
```

```
Out[9]: Unnamed: 0      0
        key           0
        fare_amount   0
        pickup_datetime 0
        pickup_longitude 0
        pickup_latitude 0
        dropoff_longitude 1
        dropoff_latitude 1
        passenger_count 0
        dtype: int64
```

Here we can see there are 14 null values in drop_off latitude and longitude. as removing 14 to 28 rows from our huge dataset will not affect our analysis so, lets remove the rows having null values

```
In [12]: 1 data.dropna(inplace=True)
        2 print(data.isnull().sum())
```


```
Unnamed: 0      0
key           0
fare_amount   0
pickup_datetime 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 0
dropoff_latitude 0
passenger_count 0
dtype: int64
```

```
In [13]: 1 import matplotlib.pyplot as plt
        2 import seaborn as sns
        3 %matplotlib inline
```

```
In [15]: 1 sns.distplot(data['fare_amount'])
```

```
c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[15]: <AxesSubplot:xlabel='fare_amount', ylabel='Density'>
```




In distribution plot also it can be seen that there are some values which are negative fare

```
In [18]: 1 sns.distplot(data['pickup_latitude'])
```

```
c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[18]: <AxesSubplot:xlabel='pickup_latitude', ylabel='Density'>
```




Here we can see minimum value is going to be less than even -3000 which is not correct value and also on positive side also going more than 2000

```
In [20]: 1 sns.distplot(data['pickup_longitude'])
```

```
c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[20]: <AxesSubplot:xlabel='pickup_longitude', ylabel='Density'>
```




Here also negative and positive values are exceeding far behind the real limit.

```
In [21]: 1 sns.distplot(data['dropoff_longitude'])
```

```
c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```


```
Out[21]: <AxesSubplot:xlabel='dropoff_longitude', ylabel='Density'>
```



```
In [23]: 1 #Similarly here also same issue
          2 sns.distplot(data['dropoff_latitude'])
```

c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[23]: <AxesSubplot:xlabel='dropoff_latitude', ylabel='Density'>
```



```
In [24]: 1 print("drop_off latitude min value",data["dropoff_latitude"].min())
          2 print("drop_off latitude max value",data["dropoff_latitude"].max())
          3 print("drop_off longitude min value", data["dropoff_longitude"].min())
          4 print("drop_off longitude max value",data["dropoff_longitude"].max())
          5 print("pickup latitude min value",data["pickup_latitude"].min())
          6 print("pickup latitude max value",data["pickup_latitude"].max())
          7 print("pickup longitude min value",data["pickup_longitude"].min())
          8 print("pickup longitude max value",data["pickup_longitude"].max())
```

```
drop_off latitude min value -881.9855130000001
drop_off latitude max value 872.6976279999999
drop_off longitude min value -3356.6663
drop_off longitude max value 1153.5726029999998
pickup latitude min value -74.01551500000001
pickup latitude max value 1644.421482
pickup longitude min value -1340.64841
pickup longitude max value 57.418457
```

we can see what is range of latitude and longitude of our test dataset, lets keep the range same in our train set so that even noisy data is remove and we have only the values which belongs to new york

```
In [25]: 1 min_longitude=-74.263242,
          2 min_latitude=40.573143,
          3 max_longitude=-72.986532,
          4 max_latitude=41.709555
```

In [26]: 1 #lets drop all the values which are not coming in above boundary, as th

In [27]: 1 tempdf=data[(data["dropoff_latitude"]<min_latitude) | (data["pickup_lat
2 print("before dropping",data.shape)
3 data.drop(tempdf.index,inplace=True)
4 print("after dropping",data.shape)

before dropping (199999, 9)
after dropping (195732, 9)

In [28]: 1 #lets remove all those rows where fare amount is negative

In [29]: 1 print("before dropping", data.shape)
2 train_df=data[data['fare_amount']>0]
3 print("after dropping", data.shape)

before dropping (195732, 9)
after dropping (195732, 9)

On different day and time there would be different price like during eveing price would be more compare to afternoon, during christmas price would be different and similarly on weekends price would be different compare to week days. so lets create some extra features which will take care of all these things

In [30]: 1 import calendar
2 data['day']=data['pickup_datetime'].apply(lambda x:x.day)
3 data['hour']=data['pickup_datetime'].apply(lambda x:x.hour)
4 data['weekday']=data['pickup_datetime'].apply(lambda x:calendar.day_name
5 data['month']=data['pickup_datetime'].apply(lambda x:x.month)
6 data['year']=data['pickup_datetime'].apply(lambda x:x.year)

In [31]: 1 data.head()

Out[31]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06+00:00	-73.999817	40.73
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56+00:00	-73.994355	40.72
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00+00:00	-74.005043	40.74
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21+00:00	-73.976124	40.75
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00+00:00	-73.925023	40.74

In []: 1 #here we can see that week are in monday , tuesday and so on. So we need

```
In [32]: 1 data.weekday = data.weekday.map({'Sunday':0, 'Monday':1, 'Tuesday':2, 'Wed
```

```
In [33]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 195732 entries, 0 to 199999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            195732 non-null  int64
1   key                   195732 non-null  object
2   fare_amount           195732 non-null  float64
3   pickup_datetime       195732 non-null  datetime64[ns, UTC]
4   pickup_longitude      195732 non-null  float64
5   pickup_latitude       195732 non-null  float64
6   dropoff_longitude     195732 non-null  float64
7   dropoff_latitude      195732 non-null  float64
8   passenger_count       195732 non-null  int64
9   day                   195732 non-null  int64
10  hour                  195732 non-null  int64
11  weekday               195732 non-null  int64
12  month                 195732 non-null  int64
13  year                  195732 non-null  int64
dtypes: datetime64[ns, UTC](1), float64(5), int64(7), object(1)
memory usage: 22.4+ MB
```

```
In [34]: 1 # we will keep only those rows where number of passangers are less than
```

```
In [35]: 1 data=data[data['passenger_count']<=8]
```

```
In [36]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 195731 entries, 0 to 199999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            195731 non-null  int64
1   key                   195731 non-null  object
2   fare_amount           195731 non-null  float64
3   pickup_datetime       195731 non-null  datetime64[ns, UTC]
4   pickup_longitude      195731 non-null  float64
5   pickup_latitude       195731 non-null  float64
6   dropoff_longitude     195731 non-null  float64
7   dropoff_latitude      195731 non-null  float64
8   passenger_count       195731 non-null  int64
9   day                   195731 non-null  int64
10  hour                  195731 non-null  int64
11  weekday               195731 non-null  int64
12  month                 195731 non-null  int64
13  year                  195731 non-null  int64
dtypes: datetime64[ns, UTC](1), float64(5), int64(7), object(1)
memory usage: 22.4+ MB
```



```
In [37]: 1 #here key column and pickup_datetime columns are not needed as we have
```

```
In [38]: 1 data.drop(["key","pickup_datetime"], axis=1, inplace=True)
```

```
In [39]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 195731 entries, 0 to 199999
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            195731 non-null  int64
1   fare_amount           195731 non-null  float64
2   pickup_longitude      195731 non-null  float64
3   pickup_latitude       195731 non-null  float64
4   dropoff_longitude     195731 non-null  float64
5   dropoff_latitude      195731 non-null  float64
6   passenger_count       195731 non-null  int64
7   day                   195731 non-null  int64
8   hour                  195731 non-null  int64
9   weekday               195731 non-null  int64
10  month                 195731 non-null  int64
11  year                  195731 non-null  int64
dtypes: float64(5), int64(7)
memory usage: 19.4 MB
```

lets divide the data set into train and validation test set

```
In [40]: 1 from sklearn.model_selection import train_test_split
```

```
In [54]: 1 x=data.drop("fare_amount", axis=1)
```

```
In [55]: 1 y=data['fare_amount']
```

```
In [56]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,rand
```

```
In [57]: 1 x_train.head()
```

```
Out[57]:
```

	Unnamed: 0	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	pas
7570	51992033	-73.991973	40.742657	-73.991358	40.750086	
155037	10241908	-73.964111	40.807957	-73.966688	40.803299	
67010	48963133	-73.987658	40.700823	-73.985670	40.770540	
155236	30446807	-73.999577	40.726656	-74.007562	40.713286	
187226	40739497	-73.983377	40.738938	-73.978432	40.745286	

In [58]: 1 x_test.head()

Out[58]:

	Unnamed: 0	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
51869	5536882	-73.953347	40.767932	-73.990867	40.751295	1
44724	35054768	-73.137393	41.366138	-73.137393	41.366138	1
47705	15258057	-74.009707	40.712480	-73.962757	40.758977	1
17345	34739111	-74.016055	40.715077	-74.008840	40.711375	1
179351	53446498	-73.950474	40.784003	-73.971086	40.748328	1

In [59]: 1 x_train.shape

Out[59]: (156584, 11)

In [60]: 1 x_test.shape

Out[60]: (39147, 11)

In [61]: 1
2 *#Lets run the model.*
3 *#As we have to build regression model, Lets start with Linear regression*

In [62]: 1 from sklearn.linear_model import LinearRegression

In [63]: 1 lrmodel=LinearRegression()
2 lrmodel.fit(x_train, y_train)

Out[63]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [64]: 1 LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

Out[64]: LinearRegression(n_jobs=1, normalize=False)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [65]: 1 predictedvalues = lrmodel.predict(x_test)

```
In [66]: 1 #Lets calculate rmse for Linear Regression model
2 from sklearn.metrics import mean_squared_error
3 lrmodelrmse = np.sqrt(mean_squared_error(predictedvalues, y_test))
4 print("RMSE value for Linear regression is", lrmodelrmse)
```

RMSE value for Linear regression is 8.363019859396488

```
In [71]: 1 #Lets see with Random Forest and calculate its rmse
2 from sklearn.ensemble import RandomForestRegressor
3 # rfrmodel = RandomForestRegressor(n_estimators=100, random_state=101)
4 rfrmodel = RandomForestRegressor(n_estimators=50, random_state=101)
```

```
In [72]: 1 rfrmodel.fit(x_train,y_train)
```

Out[72]: RandomForestRegressor(n_estimators=50, random_state=101)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [74]: 1 rfrmodel_pred= rfrmodel.predict(x_test)
```

```
In [75]: 1 rfrmodel_rmse=np.sqrt(mean_squared_error(rfrmodel_pred, y_test))
2 print("RMSE value for Random forest regression is ",rfrmodel_rmse)
```

RMSE value for Random forest regression is 3.9973617568779463

```
In [76]: 1 #RandomForest Regressor is giving good value, so we can use it as final
```