

```
In [33]: 1 import pandas as pd
2 import numpy as np
```

```
In [34]: 1 data = pd.read_csv("diabetes.csv")
2 data
```

```
Out[34]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunci
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	...	...	...	...	...	...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns



```
In [35]: 1 df = pd.DataFrame(data)
2 df.head()
```

```
Out[35]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunci
0	6	148	72	35	0	33.6	0.
1	1	85	66	29	0	26.6	0.
2	8	183	64	0	0	23.3	0.
3	1	89	66	23	94	28.1	0.
4	0	137	40	35	168	43.1	2.



```
In [36]: 1 df.isnull().sum()
```

```
Out[36]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [37]: 1 plt.figure(figsize=(12,10)) # on this line I just set the size of figure
2 p=sns.heatmap(df.corr(), annot=True,cmap ='RdYlGn') # seaborn has very
```



## Manipulating and Cleaning our dataset

```
In [38]: 1 cols_clean = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesP']
2 for i in cols_clean:
3     df[i] = df[i].replace(0, np.NaN)
4     cols_mean = int(df[i].mean(skipna=True))
5     df[i] = df[i].replace(np.NaN, cols_mean)
6 data1 = df
7 data1.head().style.highlight_max(color="lightblue").highlight_min(color="lightcoral")
```

```
Out[38]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesP
0	6	148.000000	72.000000	35.000000	155.000000	33.600000	
1	1	85.000000	66.000000	29.000000	155.000000	26.600000	
2	8	183.000000	64.000000	29.000000	155.000000	23.300000	
3	1	89.000000	66.000000	23.000000	94.000000	28.100000	
4	0	137.000000	40.000000	35.000000	168.000000	43.100000	

```
In [39]: 1 print(data1.describe())
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	121.682292	72.386719	29.108073	155.28125
std	3.369578	30.435999	12.096642	8.791221	85.02155
min	0.000000	44.000000	24.000000	7.000000	14.00000
25%	1.000000	99.750000	64.000000	25.000000	121.50000
50%	3.000000	117.000000	72.000000	29.000000	155.00000
75%	6.000000	140.250000	80.000000	32.000000	155.00000
max	17.000000	199.000000	122.000000	99.000000	846.00000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.450911	0.471876	33.240885	0.348958
std	6.875366	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
In [40]: 1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 %matplotlib inline
4
5 # graph = ['Glucose', 'Insulin', 'BMI', 'Age', 'Outcome']
6 sns.set()
7 # print(sns.pairplot(data1[graph], hue='Outcome', diag_kind='kde'))
8 print(sns.pairplot(data1[graph], hue='Outcome', diag_kind='kde'))
```

<seaborn.axisgrid.PairGrid object at 0x000002A3E0B306A0>



```
In [46]: 1 # for the purpose of simplicity and analysing the most relevant data ,
2 # Glucose , Insulin and BMI
3 # defining variables and features for the dataset for splitting
4 # q_cols = ['Glucose', 'Insulin', 'BMI', 'Outcome']
5 q_cols = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', '
6
7 df = data1[q_cols]
8 print(df.head(2))
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148.0	72.0	35.0	155.0	33.6	
1	1	85.0	66.0	29.0	155.0	26.6	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0

```

In [47]: 1  ## Let's split the data into training and testing datasets
2  # split = 0.75 # 75% train and 25% test dataset
3  # total_len = len(df)
4  # split_df = int(total_len*split)
5  # train, test = df.iloc[:split_df,0:4],df.iloc[split_df:,0:4]
6  # train_x = train[['Glucose', 'Insulin', 'BMI']]
7  # train_y = train['Outcome']
8  # test_x = test[['Glucose', 'Insulin', 'BMI']]
9  # test_y = test['Outcome']
10
11 # Split the data into training and testing datasets
12 split = 0.75 # 75% train and 25% test dataset
13 total_len = len(df)
14 split_df = int(total_len * split)
15 train, test = df.iloc[:split_df], df.iloc[split_df:]
16
17 # Select the columns specified in q_cols for training and testing
18 train_x = train[q_cols[:-1]] # Exclude the 'Outcome' column from featurization
19 train_y = train['Outcome'] # Target variable
20 test_x = test[q_cols[:-1]] # Exclude the 'Outcome' column from featurization
21 test_y = test['Outcome'] # Target variable
22

```

```

In [48]: 1  a = len(train_x)
2  b = len(test_x)
3  print(' Training data =',a,'\n','Testing data =',b,'\n','Total data length =',a+b)

Training data = 576
Testing data = 192
Total data length = 768

```


```

In [49]: 1  from sklearn.neighbors import KNeighborsClassifier
2  from sklearn import metrics
3
4  def knn(x_train, y_train, x_test, y_test,n):
5      n_range = range(1, n)
6      results = []
7      for n in n_range:
8          knn = KNeighborsClassifier(n_neighbors=n)
9          knn.fit(x_train, y_train)
10         #Predict the response for test dataset
11         predict_y = knn.predict(x_test)
12         accuracy = metrics.accuracy_score(y_test, predict_y)
13         #matrix = confusion_matrix(y_test,predict_y)
14         #seaborn_matrix = sns.heatmap(matrix, annot = True, cmap="Blues")
15         results.append(accuracy)
16     return results

```

```
In [50]: 1 n= 500
          2 output = knn(train_x,train_y,test_x,test_y,n)
          3 n_range = range(1, n)
          4 plt.plot(n_range, output)
```

Out[50]: [<matplotlib.lines.Line2D at 0x2a3ec0c6100>]



```
In [51]: 1 # best k that could optimize this model is between 100 to 200 offering
          2 # ideal k value for this dataset should be 150 give or take
```

```
In [52]: 1 from sklearn.metrics import confusion_matrix
2 from sklearn.metrics import accuracy_score, precision_score, recall_score
3 y_pred = knn(train_x, train_y, test_x, test_y, n)
4 cnf_matrix = confusion_matrix(test_y, y_pred)
```

```
-----
-
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_18924\597529570.py in <module>
      2 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, fbeta_score
      3 y_pred = knn(train_x, train_y, test_x, test_y, n)
----> 4 cnf_matrix = confusion_matrix(test_y, y_pred)

c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\sklearn\metrics\_classification.py in confusion_matrix(y_true, y_pred, labels, sample_weight, normalize)
    305     (0, 2, 1, 1)
    306     """
--> 307     y_type, y_true, y_pred = _check_targets(y_true, y_pred)
    308     if y_type not in ("binary", "multiclass"):
    309         raise ValueError("%s is not supported" % y_type)


c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\sklearn\metrics\_classification.py in _check_targets(y_true, y_pred)
     82     y_pred : array or indicator matrix
     83     """
---> 84     check_consistent_length(y_true, y_pred)
     85     type_true = type_of_target(y_true, input_name="y_true")
     86     type_pred = type_of_target(y_pred, input_name="y_pred")

c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\sklearn\utils\validation.py in check_consistent_length(*arrays)
    385     uniques = np.unique(lengths)
    386     if len(uniques) > 1:
--> 387         raise ValueError(
    388             "Found input variables with inconsistent numbers of samples: %r"
    389             % [int(l) for l in lengths])

ValueError: Found input variables with inconsistent numbers of samples: [192, 499]
```

```
In [53]: 1 p = sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fn
2 plt.title('Confusion matrix', y=1.1)
3 plt.ylabel('Actual label')
4 plt.xlabel('Predicted label')
```

Out[53]: Text(0.5, 12.5, 'Predicted label')



```
In [54]: 1 # Define your KNN function to return predictions
2 def knn2(x_train, y_train, x_test, y_test, n):
3     knn = KNeighborsClassifier(n_neighbors=n)
4     knn.fit(x_train, y_train)
5     # Predict the response for the test dataset
6     predict_y = knn.predict(x_test)
7     return predict_y
```

```
In [55]: 1 n = 500
2 y_pred = knn2(train_x, train_y, test_x, test_y, n)
3 cnf_matrix = confusion_matrix(test_y, y_pred)
```

```
In [56]: 1 # Now you can calculate other metrics like accuracy, precision, recall,
2 accuracy = accuracy_score(test_y, y_pred)
3 precision = precision_score(test_y, y_pred)
4 recall = recall_score(test_y, y_pred)
5 f1 = f1_score(test_y, y_pred)
6 fbeta = fbeta_score(test_y, y_pred, beta=0.5)
```

c:\users\kedar\appdata\local\programs\python\python39\lib\site-packages\sklearn\metrics\\_classification.py:1327: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero\_division` parameter to control this behavior.  
\_warn\_prf(average, modifier, msg\_start, len(result))



```
In [57]: 1 # Print the confusion matrix and other metrics
2 print("Confusion Matrix:\n", cnf_matrix)
3 print("Accuracy:", accuracy)
4 print("Precision:", precision)
5 print("Recall:", recall)
6 print("F1 Score:", f1)
7 print("F-beta Score:", fbeta)
```

Confusion Matrix:

```
[[122  0]
 [ 70  0]]
```

Accuracy: 0.6354166666666666

Precision: 0.0


Recall: 0.0

F1 Score: 0.0

F-beta Score: 0.0

```
In [58]: 1 p = sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu", fn
2 plt.title('Confusion matrix', y=1.1)
3 plt.ylabel('Actual label')
4 plt.xlabel('Predicted label')
```

Out[58]: Text(0.5, 12.5, 'Predicted label')



```
In [60]: 1 from sklearn.model_selection import train_test_split
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
5
```

In [61]:

```
1  # Load your dataset
2  # Replace 'your_dataset.csv' with the actual file path to your dataset
3  df = pd.read_csv('diabetes.csv')
4
5  # Define your feature columns and target column
6  q_cols = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', '
7  target_col = 'Outcome'
8
9
10 # Split the data into features (X) and target (y)
11 X = df[q_cols]
12 y = df[target_col]
13
14 # Split the data into training and testing datasets
15 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
16
17 # Perform feature scaling (standardization) on the features
18 scaler = StandardScaler()
19 X_train_scaled = scaler.fit_transform(X_train)
20 X_test_scaled = scaler.transform(X_test)
21
22 # Create and train a K-nearest neighbors (KNN) classifier
23 k = 5 # You can adjust the value of k
24 knn_classifier = KNeighborsClassifier(n_neighbors=k)
25 knn_classifier.fit(X_train_scaled, y_train)
26
27 # Make predictions on the test data
28 y_pred = knn_classifier.predict(X_test_scaled)
29
30 # Evaluate the model
31 accuracy = accuracy_score(y_test, y_pred)
32 conf_matrix = confusion_matrix(y_test, y_pred)
33 classification_rep = classification_report(y_test, y_pred)
34
35 # Print the results
36 print(f"Accuracy: {accuracy}")
37 print("Confusion Matrix:")
38 print(conf_matrix)
39 print("Classification Report:")
40 print(classification_rep)
```

Accuracy: 0.6822916666666666

Confusion Matrix:

[[94 29]


[32 37]]

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.76	0.76	123
1	0.56	0.54	0.55	69
accuracy			0.68	192
macro avg	0.65	0.65	0.65	192
weighted avg	0.68	0.68	0.68	192

```
In [63]: 1 p = sns.heatmap(pd.DataFrame(conf_matrix), annot=True, cmap="YlGnBu", f
2 plt.title('Confusion matrix', y=1.1)
3 plt.ylabel('Actual label')
4 plt.xlabel('Predicted label')
```

Out[63]: Text(0.5, 12.5, 'Predicted label')



```
In [ ]: 1
```